

# DONALD B. HAVERY

Applied AI Engineer, Forward-Deployed Focus · Audit, Evals, Deployment for Production AI Systems · Agentic Orchestration · RAG & Context Engineering · Eval Gates · Security Controls · MCP · Local-First

Portland, OR Metro Area (open to Seattle / SF / NYC, remote-first, 25% travel) | [dbhavery@gmail.com](mailto:dbhavery@gmail.com) | [github.com/dbhavery](https://github.com/dbhavery) | [linkedin.com/in/dbhavery](https://linkedin.com/in/dbhavery) | [dbhavery.dev](https://dbhavery.dev)

## SUMMARY

I turn commodity AI into working software inside a real business: audit the workflow, prove it with evals, deploy it with an audit trail people can trust. Proof below is public, cloneable, and CI-green, not claimed, and one of them runs my own business and takes live payments. You hire me and the team of AI agents I build and run. They come with me.

## REPRESENTATIVE PROJECTS

### Production AI agent, live on a revenue-taking business

Live and answering customers · every price comes from a tool call against the price book, never from the model · a stream guard kills un-tooled figures before the customer sees them and escalates to a human · live Stripe checkout · built and run solo

### Policy-control evidence review

Audit stage · per-control verdicts against matched evidence, append-only audit log · MIT

### LLM & RAG evaluation release gate

Evals stage · exits non-zero when a case regresses against baseline · MIT

### Agent orchestration reliability harness

Deploy stage · unsafe-action gate, six-class failure taxonomy · MIT

### LLM reliability gateway

Cross-provider failover, from-scratch HNSW · MIT

### Also public

Local-first desktop AI assistant with a non-bypassable policy gate · self-hosted RAG file library · unattended research-to-publish pipeline · [self-correcting text-to-CAD loop](#) · situational-awareness globe · a library of reusable Claude Code skills · [github.com/dbhavery](https://github.com/dbhavery)

## TECHNICAL SKILLS

|                     |                                                                                                              |
|---------------------|--------------------------------------------------------------------------------------------------------------|
| Agentic & Tool Use  | Multi-agent orchestration, intent routing, autonomous loops, MCP, fallback chains                            |
| RAG & Retrieval     | Hybrid retrieval (BM25 + dense), ChromaDB, from-scratch HNSW, context engineering                            |
| Evals & Reliability | Deterministic offline assertions, regression detection, CI exit-code gating, circuit breakers, retry budgets |
| Claude Code         | CLAUDE.md project context, hooks, skills, subagents, plan-first workflow                                     |
| LLM Stack           | Claude, Gemini, Qwen, Ollama local serving; provider-agnostic routing; LoRA / QLoRA fine-tuning              |
| Backend & Product   | Python (FastAPI), Rust (Tauri), Node.js, Postgres / SQLite / Redis, Docker; Next.js, React 19, TypeScript    |

## EXPERIENCE

### Applied AI Engineer, Forward-Deployed Focus, Independent

Current

- Audit, prove, deploy: proof below, cloneable and tested, not claimed
- Daily Claude Code work: CLAUDE.md context, hooks, plan before build, results verified by running them
- Portfolio at [dbhavery.dev](https://dbhavery.dev)

### Other Experience

Founder / owner-operator, built and run solo alongside full-time work: [angrynirds](#), a web and AI studio for small business, live and taking payments, with the production agent above on the front desk; [Trucking Tutor](#), a live CDL training site; [role-based fuel-delivery operations platform](#); DTC premium grooming brand; independent investigative case work. 30 years of line-haul freight under DOT and FMCSA rules, Class A CDL with hazmat, tanker, and triples, and a company-certified driver trainer, certified through the carriers I worked for.

## RESEARCH

“**Reference-Relative Believability**”, independent research paper introducing a reference-relative method for evaluating the perceived believability of AI-generated media, archived on Zenodo. [DOI 10.5281/zenodo.21138188](https://doi.org/10.5281/zenodo.21138188)

## EDUCATION

Capella University, B.S. Computer Science.

## CERTIFICATIONS

**Anthropic, Teaching the AI Fluency Framework** (April 2026), delivered with University College Cork, Ringling College of Art + Design, and Ireland’s Higher Education Authority and National Forum, plus **14 further Anthropic Academy certifications** across Claude, Claude Code, MCP, and enterprise deployment on Bedrock and Vertex AI

**IBM, AI Product Manager Professional Certificate**, ten courses | **DeepLearning.AI, Generative AI with LLMs**